

Piece Square Tables Factorisation

Abstract

This document describes a method of lossy Piece Square Table (PSQT) compression such that the result uses fewer *tokens* than lossless compression methods while maintaining engine's ability to play at a competitive level. The compression is based on matrix factorisation obtained using genetic algorithm.

1 Introduction

The main idea behind this method stems from the fact that piece square tables share some similarities independent of the piece type and game phase they are tied to. For example pawns are given bonus for reaching 6th and 7th rank regardless of the phase the game is in. During endgame that bonus is greater but we are mainly interested in the fact that it exists and is applied to the same squares. Similarly pieces are often penalised for being placed on the edges of the board.

These similarities can be described using binary *components* which in turn can be represented using bitboards. For example a bitboard with only the bits corresponding to the 6th and 7th rank set represents a component that can be used to either reward or penalise these squares.

The idea is to perform an element-wise linear combination of these components to obtain an approximation of a given table. While the components are shared among all tables/pieces the weights used in the linear combination are piece-specific.

The challenge is to find the optimal structure of components, values of weights and the number of both so as to minimise the number of literals needed to store the data as well as to minimise the Elo loss of an engine using that approximation.

2 Method

2.1 Formal Description

To formalise the ideas introduced in the previous section I introduce the following definitions

- \mathcal{P} – the set of target piece square tables (matrices of size 8 by 8)
- $P \in \mathbb{Z}^{64 \times 12}$ – matrix representing the tables from \mathcal{P} where each column $P_{:,i}$ is a matrix flattened to a column vector representing the i – th table
- $q \in \mathbb{Z}$ – piece value of the i -th piece.
- $c_i \in \{0, 1\}^{64}$ – the i -th component represented using a binary vector of size 64.
- $C \in \{0, 1\}^{64 \times n}$ – binary component matrix such that $C_{:,i} = c_i$ where n is the number of components
- $w_i \in \mathbb{Z}^n$ – the i -th vector of weights corresponding.
- $W \in \mathbb{Z}^{n \times 12}$ – weight matrix such that $W_{:,i} = w_i$.

Using these definitions we can write the approximation of P as matrix multiplication of C and W

$$P \approx \hat{P} = CW \tag{1}$$

Our goal is to solve the following task

$$\arg \min_{C, W} \left\| P - \hat{P} \right\|_F^2 \tag{2}$$

where $\|\cdot\|_F$ is Frobenius norm. I chose this loss metric because it's simple, intuitive and has good mathematical properties.

2.2 Solving

2.2.1 Matrix Factorisation with Binary Components

At first I've tried finding the matrices C and W using the algorithm described in [1]. It seemed perfect for the problem but I found that because of the limitations of the algorithm it could be outperformed by my own solution which is described in the following section.

2.2.2 Genetic Algorithm

To solve the minimisation problem I've used *genetic algorithm* which simulates generations of specimens to imitate evolution using a specified fitness function to select best specimens.

Each specimen contains *genes* composed of matrices \tilde{C} and \tilde{W} which are contenders to become C and W . During every generation each specimen s_i is evaluated using a fitness function

$$f(s_i) = - \left\| P - \tilde{C}_i \tilde{W}_i \right\|_F^2 \quad (3)$$

The minus sign represents the fact that evolution algorithms in general try to *maximise* the fitness function. The greater the fitness function value the better is the specimen.

A new generation is created by maintaining top k specimens from the previous generation (to keep the current best solution) and creating a set of new specimens. Each new specimen is created by selecting a parent (a specimen from the previous generation) at random (favouring best performing specimens) and slightly mutating it.

Mutation consists of flipping an element inside \tilde{C} from 0 to 1 or vice versa and/or increasing/decreasing the value of a random element inside \tilde{W} by a small amount.

After many generations the specimens eventually converge to a local minimum of the fitness function. At that moment \tilde{C} and \tilde{W} are selected as C and W

2.3 Additional Considerations

2.3.1 Importance Mask

One way we can improve the solutions is to realise that some values inside P are not important in any way. Specifically the values of the 1st and 8th rank for pawns in no way influence the evaluation because pawns can never be there (under standard chess rules). We can represent this fact but creating an importance mask $M \in \{0, 1\}^{64 \times 12}$ which consists entirely of ones except for elements corresponding to the aforementioned squares. We can then rewrite the fitness function as

$$f(s_i) = - \left\| M \odot (P - \tilde{C}_i \tilde{W}_i) \right\|_F^2 \quad (4)$$

where \odot represents element wise multiplication.

2.3.2 Iterative Component Generation

To ensure good convergence the mutation rate inside a genetic algorithm should be kept low. However in our case it can lead to several component-weight pairs to evolve similarly. This results in optimal solutions because simply removing one of the identical components and summing the weight vectors would result in the same approximation with less data.

To address this problem I perform the genetic algorithm in several iterations. At iteration $i \in [n]$ specimens only have i components and $12 \cdot i$ weights to optimise. After the evolution has converged we can advance to the next iteration by adding a new component consisting of only zeros and a new set of weights also consisting of only zeros.

In practice this approach, while slower, leads to much better results.

2.3.3 Mean Adjusted PSQTs

Another thing I've implemented is mean adjusted PSQTs values. For each table $p \in \mathcal{P}$ and corresponding piece value q I calculate

$$\bar{p} = \lfloor \frac{\sum_{i,j=1}^{i,j=8} p_{ij}}{64} \rfloor \quad (5)$$

and update the values

$$\begin{aligned} p &\leftarrow p - \bar{p} \\ q &\leftarrow q + \bar{p} \end{aligned}$$

The logic behind this method is that we do not want factorise the bias of a given table.

I have not tested this approach thoroughly but it seems to lead to slightly better results.

2.4 Optimal n

The last thing to solve is to find the optimal n i.e. the optimal number of components so that we save tokens by using few literals but get a good approximation of P at the same time.

I have no good theoretical approach to this problem but from many rounds testing I've found that $n = 12$ is probably the best value. For values less than 12 the Elo

loss was too big to be justified and for values greater than 12 I've encountered quickly diminishing returns. However there is a good change that a greater value could result in a much better approximation. As such this problem remains unsolved.

2.5 Packing/Unpacking

Finally there is the issue of packing and unpacking the data. Apart from C and W we also need to include q to fully reconstruct the value of any given piece on any given square.

We can leverage the fact that the chess board has 64 squares so each component can be represented as a bitboard i.e. single *ulong*. Similarly a *decimal* can store 12 useful bytes of data. This means that if we limit our weights to only take up 8 bits of memory we can use one decimal to store the weights corresponding to a given component. In practice the weights are in range $[-128, 127]$.

To store q we can divide the values by $d \in \mathbb{N}$ so that they fit in the specified weight range. Then we can encode them as another set of weights with a component that is a matrix of size $64 \times n$ and contains only d . Since for $d > 1$ we cannot represent this component using bitboards we need to declare it implicitly. I achieve this by using a ternary operator that checks if the index of a given weight set is $n + 1$ and multiply the weights by d otherwise I multiply the weight by the corresponding bit inside the component.

The code snippet for unpacking is as follows

```
UnpackedPestoTables = Enumerable.Range(0, 768).Select(i =>
    weights.Select((w, j) =>
        (sbyte)Buffer.GetByte(decimal.GetBits(w), i / 64) * (j == n ? d :
            (int)((components[j] >> i % 64) % 2))
    ).Sum()).ToArray();
```

where `weights` and `components` are arrays of respectively decimals and ulongs of size n .

3 Results

Using $n = 12$ (25 literals) I was able to find matrices C and W such that the engine using approximated PeSTO PSQTs was comparable to the engine using lossless compression.




Figure 1: PeSTO Tables




Figure 2: Approximated PeSTO Tables

Running an SPRT between the two engines with bound $[-3,0]$ resulted in a reported -20.4 ± 8.6 Elo difference. Overall the method is 29 tokens less expensive than lossless compression methods. I consider this a success since my goal was to find an approximation which would result in ratio of Elo lost to tokens saved being less than 1.

References

- [1] Martin Slawski, Matthias Hein, and Pavlo Lutsik. Matrix factorization with binary components, 2014.